

Block Drivers

What to Expect?

- ★ Why the need for the Block Layer?
- ★ Decoding a Block Device in Linux
- ★ Role of Block Drivers
- ★ Writing a Block Driver

Block vs Character

★ Concept Similarities

- Device Registration
- Usage of Device Files
- Major & Minor number
- File Operations


★ Then, why a different category?

- To access block-oriented devices (Really?)
- To achieve buffering for efficiency
- To enable a generic caching abstraction
- To provide a device independent block access of data


System-wide Block Devices

- ★ Category is to Major
 - IDE: 3; SCSI: 8; ...
- ★ Naming Convention (Try: `ls -l /dev/sd*`)
 - IDE: `hd*`; SCSI: `sd*`; ...
- ★ Disk is to Minor
 - Typically limited to 4 per system, represented using a, b, ...
- ★ Partition also is to Minor
 - $256 / 4 = 64$ to each disk
 - First one for the whole disk, e.g. `hda`
 - Remaining for the partitions, thus limiting to 63, e.g. `hda1`

The Generic Hard Disk



The Generic Hard Disk



Computing a Generic Hard Disk

★ Example (Hard Disk)

- ▶ Heads (or Platters): 0 – 9
- ▶ Tracks (or Cylinders): 0 – 24
- ▶ Sectors: 1 – 64

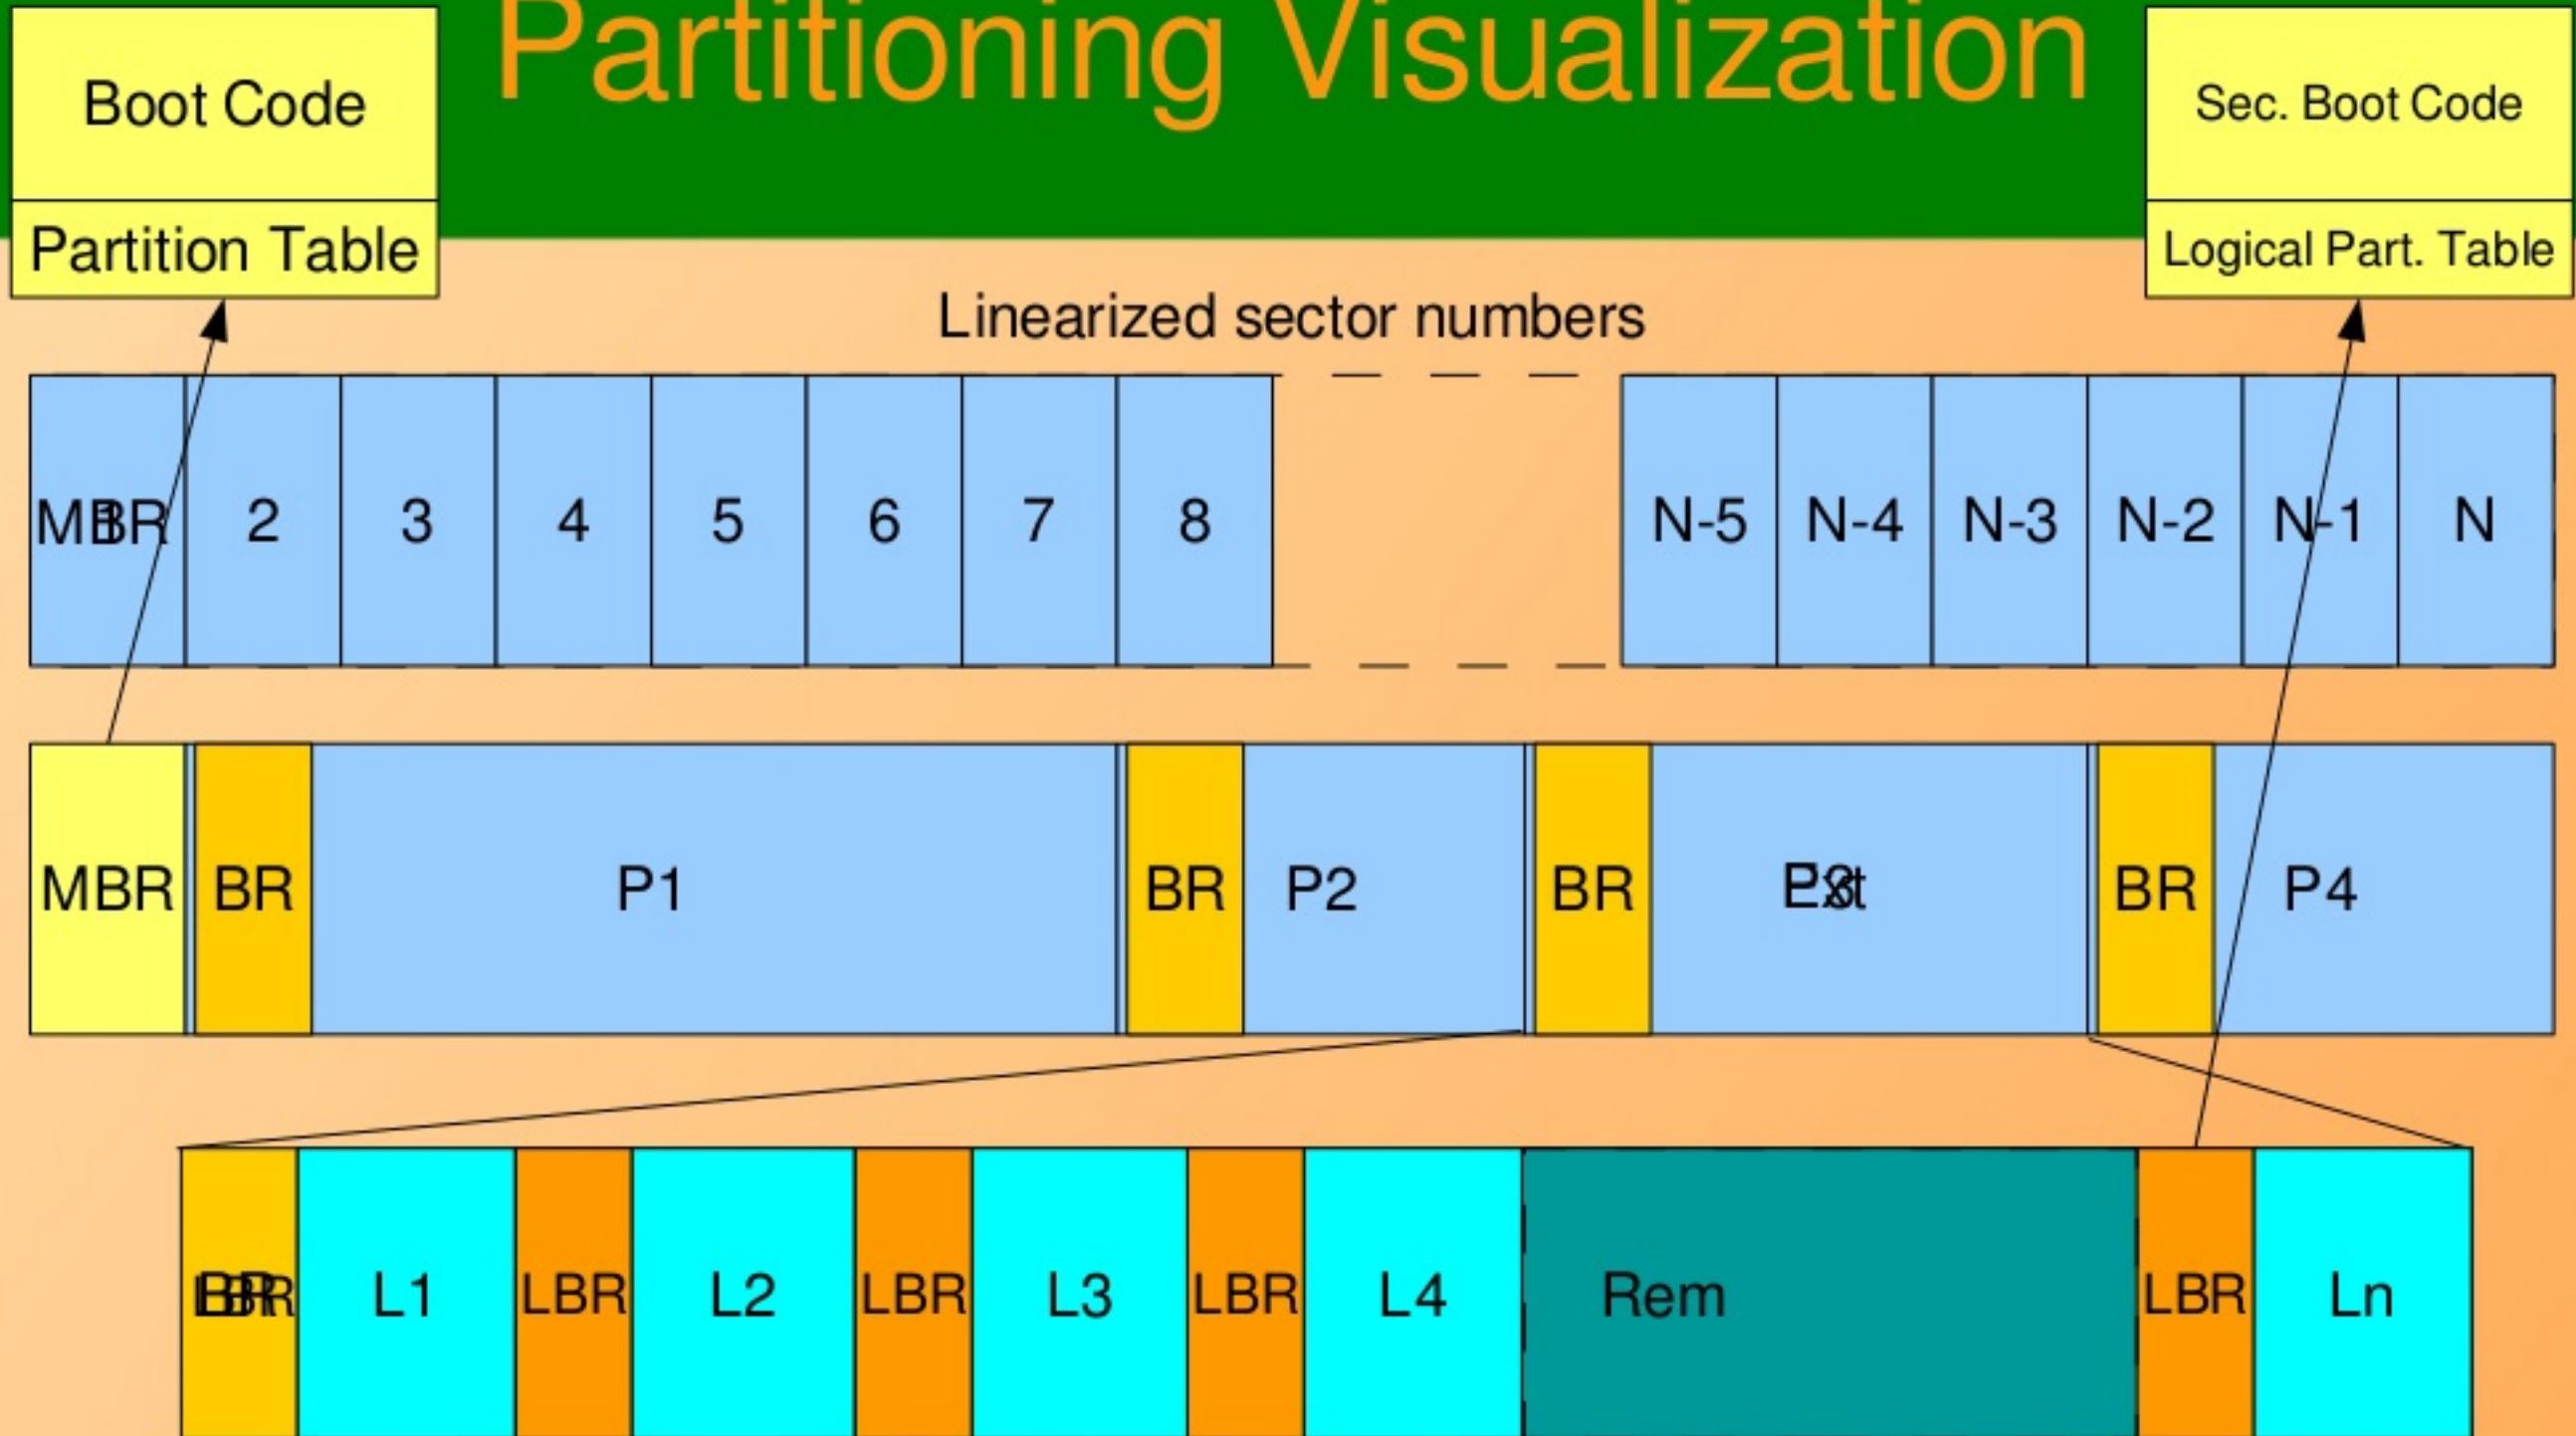
★ Size of the Hard Disk

- ▶ $10 \times 25 \times 64 \times 512 \text{ bytes} = 8000\text{KiB}$

★ Device independent numbering

- ▶ $(h, t, s) \rightarrow 64 * (10 * t + h) + s \rightarrow (1 - 16000)$

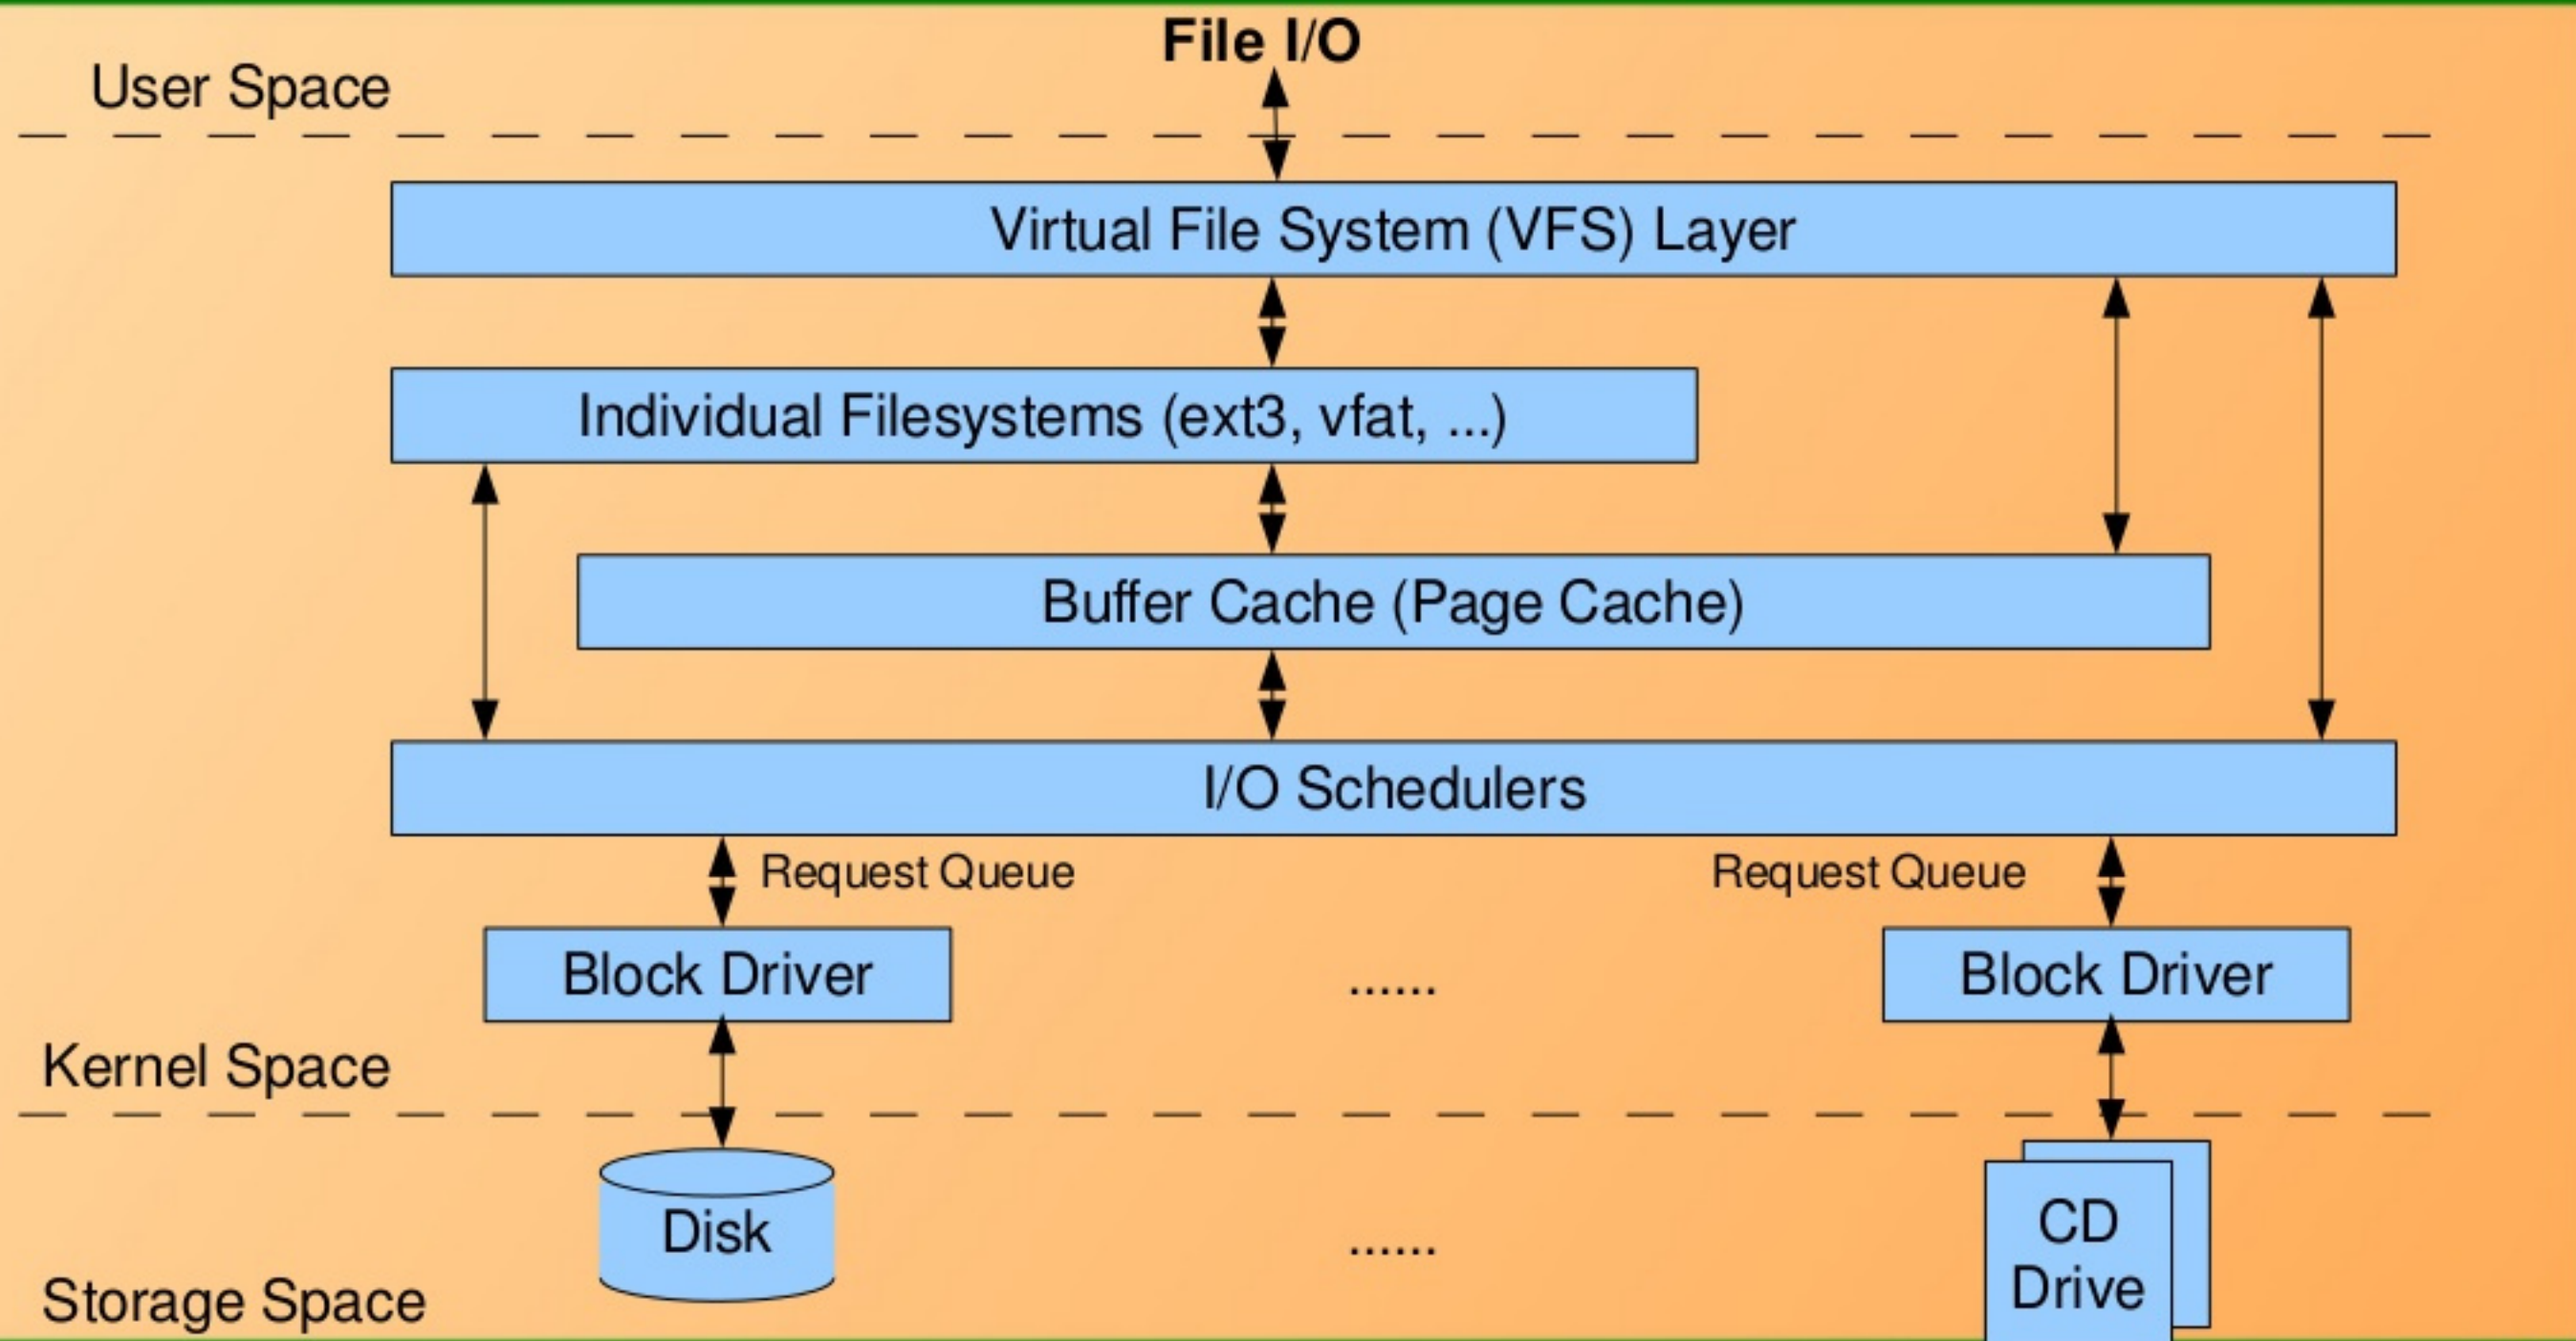
Partitioning Visualization



Partitioning a Block Device

- ★ First Sector – Master Boot Record (MBR)
 - Contains Boot Info
 - Contains Physical Partition Table
- ★ Maximum Physical Partitions: 4
 - At max 1 as Extended Partition
 - Rest as Primary Partition
- ★ Extended could be further partitioned into
 - Logical Partitions
- ★ In each partition
 - First Sector – Boot Record (BR)
 - Remaining for File System / Format
 - Extended Partition BR contains the Logical Partition Table

Block Input / Output



Now, let's write a Driver
to
Achieve the Purpose

Block Registration

★ Driver Registration

- Header: `<linux/fs.h>`
- APIs
 - `int register_blkdev(major, name);`
 - `int unregister_blkdev(major, name);`

★ Disk Drive Registration

- Header: `<linux/genhd.h>`
- Data Structure: `struct gendisk *gd`
- APIs
 - `struct gendisk *alloc_disk(minors); void put_disk(gd);`
 - `void add_disk(gd); void del_gendisk(gd);`

struct gendisk

- ★ int major
- ★ int first_minor
- ★ int minors
- ★ char disk_name[32]
- ★ struct block_device_operations *fops
- ★ struct request_queue *queue
- ★ int flags (GENHD_FL_REMOVABLE, ...)
- ★ sector_t nr_sects → struct hd_struct part0
- ★ void *private_data

struct hd_struct

- ★ sector_t start_sect
- ★ sector_t nr_sects
- ★ sector_t alignment_offset
- ★ ...

Block Device Operations

- ★ Header: `<linux/blkdev.h>`
- ★ System Calls (till `<= 2.6.27`)
 - `int open(struct inode *i, struct file *f);`
 - `int close(struct inode *i, struct file *f);`
 - `int ioctl(struct inode *i, struct file *f, cmd, arg);`
 - `int media_changed(struct gendisk *gd);`
 - `int revalidate_disk(struct gendisk *gd);`
 - ...
- ★ Other Important Fields
 - `struct module *owner;`

Block Device Operations

★ Header: <linux/blkdev.h>

★ System Calls (after 2.6.27)

- `int (*open)(struct block_device *, fmode_t);`
- `int (*release)(struct block_device *, fmode_t);`
- `int (*ioctl)(struct block_device *, fmode_t, cmd, arg);`
- `int (*media_changed)(struct gendisk *gd);`
- `int (*revalidate_disk)(struct gendisk *gd);`
- `int (*getgeo)(struct block_device *, struct hd_geometry *);`
- ...

★ Other Important Fields

- `struct module *owner;`

Request Queues & Processing

★ Header: `<linux/blkdev.h>`

★ Types

- `request_queue_t *q;`
- `request_fn_proc rqf;`
- `struct request *req;`

★ APIs

- `q = blk_init_queue(rqf, lock);`
- `blk_cleanup_queue(q);`
- `req = blk_fetch_request(q);`

Requests

★ Interfaces

- ▶ `rq_data_dir(req)` – Operation type
 - zero: read from device
 - non-zero: write to the device
- ▶ `blk_req_pos(req)` – Starting sector
- ▶ `blk_req_sectors(req)` – Total sectors
- ▶ Iterator for extracting buffers from `bio_vec`

★ Request Function

- ▶ `typedef void (*request_fn_proc)(request_queue_t *queue);`

Disk on RAM

★ Let's try out the RAM Block Driver

- Horizontal: Disk on RAM
 - `ram_device.c`, `ram_device.h`
- Vertical: Block Driver
 - `ram_block.c`

★ Useful commands

- `blockdev`
- `dd`
- `fdisk`

What all have we learnt?

- ★ Understood the need for the Block Layer
- ★ Decoding a Block Device in Linux
- ★ Role of Block Drivers
- ★ Writing a Block Driver
 - Registration
 - Block Device Operations
 - Request & Request Queues

Any Queries?