

Interrupts

What to Expect?

- ★ W's & How's of Interrupts?
- ★ The IRQ
- ★ Interrupt Handling & Handlers
- ★ Soft IRQs

W's of an Interrupt?

★ What is an Interrupt?

- Intervention to get Attention
- Here, from the Devices to the CPU

★ Why is an Interrupt required?

- Mismatch of Device & CPU speeds
- Get overall better efficiency & latency

How does interrupt work?

- ★ Interrupt is a Hardware thing
- ★ It needs Physical connection
- ★ In Processors
 - Typically have one or two interrupt lines
 - Handled by a Interrupt Controller, called PIC
 - And PIC informs the Processor through its one line
- ★ In Micro-controllers
 - Each GPIO can act as an interrupt line
 - Interrupt Controller is in-built
- ★ In both cases, the CPU then decodes the IRQ
- ★ And calls the corresponding registered interrupt handler

Interrupt ReQuest (IRQ)

- ★ Number derived by CPU & Board
- ★ Programmed or Hard Coded with Interrupt Controller
- ★ Hence, one key component of LSPs
- ★ In x86
 - ▶ CPU Specific: 0x00 to 0x1F (Hard coded)
 - ▶ Board Specific: 0x20 to 0xFF
 - Header: `<asm/interrupt.h>` → `<asm/irq.h>` → `irq*.h`
- ★ In Microcontrollers
 - ▶ Depends on the controller
 - Header: `<asm/interrupt.h>` → `<asm/irq.h>` → `irq*.h`

Programming Interface

★ Header: `<asm/interrupt.h>`

★ Type:

- ▶ `typedef irqreturn_t (*irq_handler_t)(int, void *);`

★ APIs

- ▶ `int request_irq(unsigned int irq, irq_handler_t handler, unsigned long flags, const char *name, void *dev_id);`
- ▶ `void free_irq(unsigned int irq, void *dev_id);`
- ▶ `int can_request_irq(irq, flags);`

★ Flags

- ▶ `IRQF_TRIGGER_RISING, ..., IRQF_TRIGGER_HIGH, ...`
- ▶ `IRQF_SAMPLE_RANDOM`
- ▶ `IRQF_SHARED, ...`

IRQ Handler Do's & Don'ts

- ★ No sleeping – direct or indirect
 - ▶ Example: `schedule_timeout`, `input_register_device`
- ★ No mutexes
 - ▶ Rather use `spin_locks`, if you must
- ★ Can't exchange data with User Space
 - ▶ In their own context. `current` is invalid
- ★ If lot to be done, break it into two
 - ▶ Top Half & Bottom Half
- ★ Need not be re-entrant
- ★ May get interrupted by higher priority interrupt handlers

The Additional Info

★ IRQ Control

- `enable_irq(irq);`
- `disable_irq(irq);`

★ IRQ Handler returns

- `IRQ_NONE`
- `IRQ_HANDLED`

★ Check for execution state

- `in_interrupt();`

★ Synchronous interrupts, treated alike

- Exceptions to report grave runtime errors
- Software interrupts such as the `int 0x80` used for system calls in x86 architecture

Message Signaled Interrupts

- ★ Specifically for PCI Devices
- ★ Advantages
 - No sharing, No sync issues, More interrupts
- ★ Modes: MSI or MSI-X (only one at a time)
- ★ MSI (since PCI 2.2)
 - Special Address: PCI config space
 - Interrupts / Device: Upto 32 in powers of 2
- ★ MSI-X (since PCI 3.0)
 - Special Address: Bus address
 - Interrupts / Device: Sparse & upto 2048

MSI/MSI-X APIs

★ MSI

- ▶ `int pci_enable_msi(struct pci_dev *);`
- ▶ `int pci_enable_msi_block(struct pci_dev *, int cnt);`
- ▶ `void pci_disable_msi(struct pci_dev *);`

★ MSI-X

- ▶ `int pci_enable_msix(struct pci_dev *, struct msix_entry *, int nvec);`
 - `struct msix_entry`
`{ u16 vector /* allocated irq */, entry; };`
- ▶ `void pci_disable_msix(struct pci_dev *);`

Soft IRQs

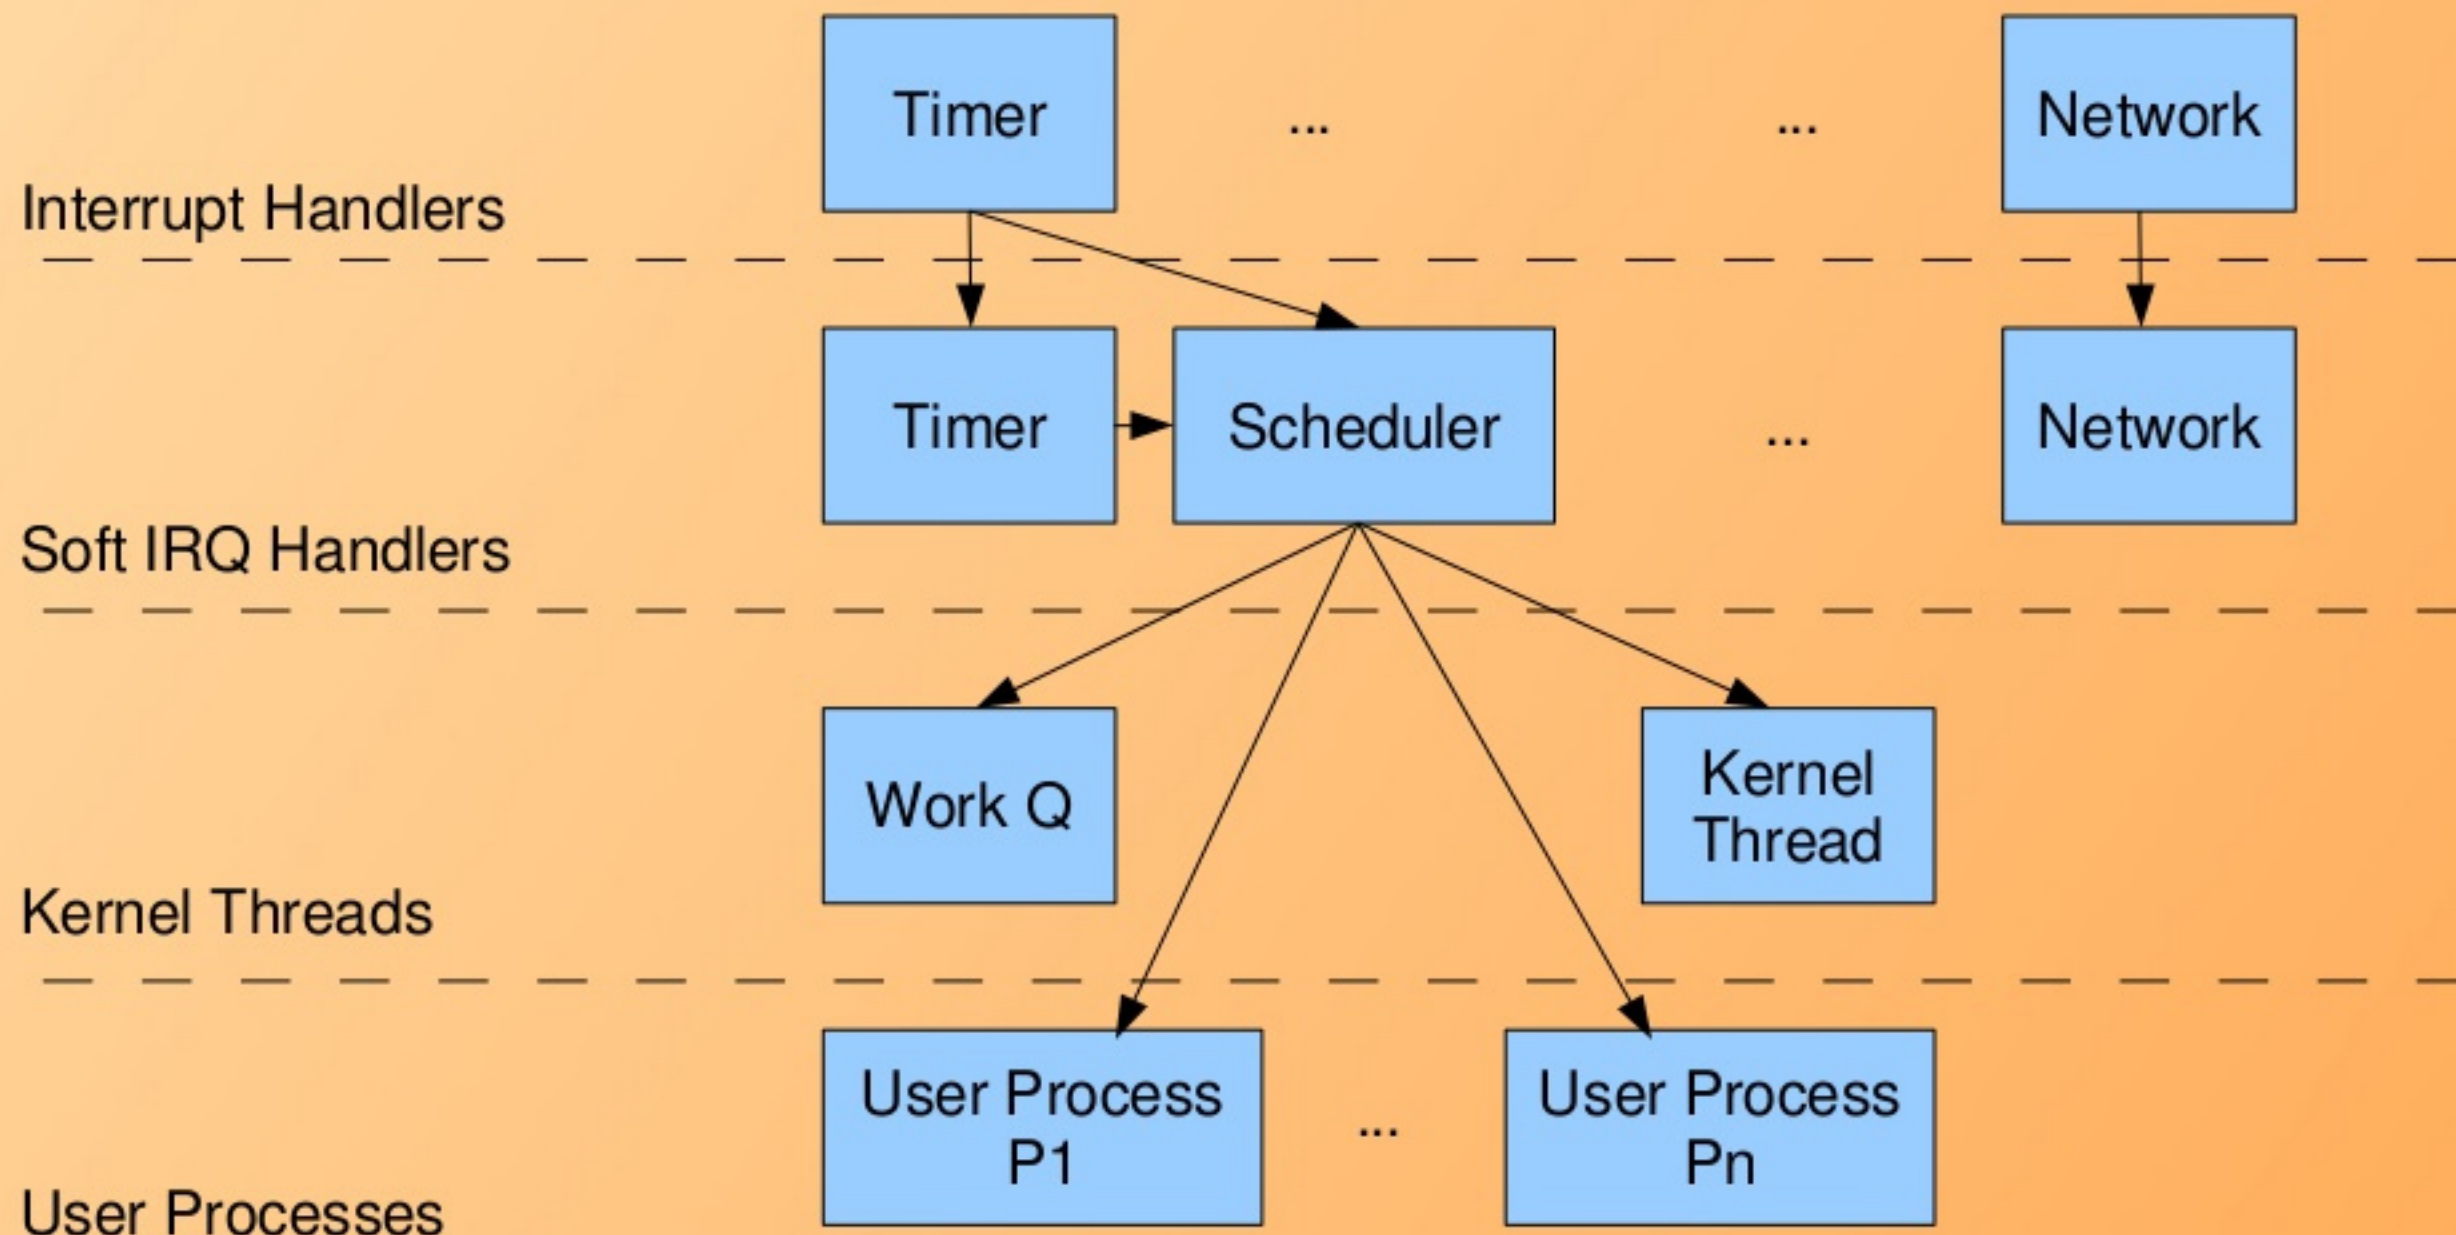
What is a Soft IRQ?

- ★ Basic Bottom Half
- ★ Synchronous Interrupt
- ★ Have strong locking requirements
- ★ Used for the performance sensitive subsystems
- ★ Not same as Software Interrupt

Typical Soft IRQs

- ★ Timer
- ★ Network
- ★ Block
- ★ Tasklet
- ★ Scheduler
- ★ High Resolution Timer

Linux Execution Model



Programming

★ Header: `<linux/interrupt.h>`

★ APIs

- ▶ `void open_softirq(int nr, void (*action)(struct softirq_action *));`
- ▶ `void raise_softirq(unsigned int nr);`

Top & Bottom Halves

★ Top Half

- ▶ Registered through request_irq

★ Bottom Half

▶ Tasklet

- Soft IRQ Context
- Fast & Atomic
- Only different tasklets can run on different CPUs

▶ Work Queue

- Special Kernel Process Context
- Allowed to sleep
- Can run simultaneously on different CPUs

What all have we learnt?

- ★ W's & How's of Interrupts?
- ★ The IRQ
- ★ Interrupt Handling & Handlers
- ★ Message Signalled Interrupts
- ★ Soft IRQs: The Bottom Halves

Any Queries?