

Kernel Programming

What to Expect?

- ☆ How to do programming in “Kernel C” for
 - Achieving Concurrency
 - Keeping Time
 - Providing Delays
 - Timer Control

Concurrency

Concurrency with Locking

★ Mutexes

- Header: `<linux/mutex.h>`
- Type: `struct mutex`
- APIs
 - `DEFINE_MUTEX`
 - `mutex_is_locked`
 - `mutex_lock`, `mutex_trylock`, `mutex_unlock`

★ Semaphores

- Header: `<linux/semaphore.h>`
- Type: `struct semaphore`
- APIs
 - `sema_init`
 - `down`, `down_trylock`, `down_interruptible`, `up`

Concurrency w/ Locking (cont.)

★ Spin Locks

- ▶ Header <linux/spinlock.h>
- ▶ Type: `spinlock_t`
- ▶ APIs
 - `spin_lock_init`
 - `spin_[try]lock`, `spin_unlock`

★ Reader-Writer Locks

- ▶ Header: <linux/spinlock.h>
- ▶ Type: `rwlock_t`
- ▶ APIs
 - `read_lock`, `read_unlock`
 - `write_lock`, `write_unlock`

Concurrency without Locking

★ Atomic Variables

- ▶ Header: `<asm-generic/atomic.h>`
- ▶ Type: `atomic_t`
- ▶ Macros
 - `ATOMIC_INIT`
 - `atomic_read`, `atomic_set`
 - `atomic_add`, `atomic_sub`, `atomic_inc`, `atomic_dec`
 - `atomic_xchg`

Concurrency w/o Locking (cont.)

★ Atomic Bit Operations

- ◆ Header: `<linux/bitops.h>`

- ◆ APIs

- `rol8, rol16, rol32, ror8, ror16, ror32`
- `find_first_bit, find_first_zero_bit`
- `find_last_bit`
- `find_next_bit, find_next_zero_bit`

- ◆ Header: `<asm-generic/bitops.h>`

- ◆ APIs

- `set_bit, clear_bit, change_bit`
- `test_and_set_bit, test_and_clear_bit, test_and_change_bit`

Wait Queues

★ Wait Queues

- ▶ Header: `<linux/wait.h>`
- ▶ Wait Queue Head APIs
 - `DECLARE_WAIT_QUEUE_HEAD(wq);`
 - `wait_event_interruptible(wq, cond);`
 - `wait_event_interruptible_timeout(wq, cond, timeout);`
 - `wake_up_interruptible(&wq);`
 - ... (non-interruptible set)
- ▶ Wait Queue APIs
 - `DECLARE_WAITQUEUE(w, current);`
 - `add_wait_queue(&wq, &w);`
 - `remove_wait_queue(&wq, &w);`

Time Keeping

Time since Bootup

- ★ tick – Kernel's unit of time. Also called jiffy
- ★ HZ – ticks per second
 - Defined in Header: `<linux/param.h>`
 - Typically, 1000 for desktops, 100 for embedded systems
- ★ 1 tick = 1ms (desktop), 10ms (embedded systems)
- ★ Variables: `jiffies` & `jiffies_64`
 - Header: `<linux/jiffies.h>`
 - APIs
 - `time_after`, `time_before`, `time_in_range`, ...
 - `get_jiffies_64`, ...
 - `msec_to_jiffies`, `timespec_to_jiffies`, `timeval_to_jiffies`, ...
 - `jiffies_to_msec`, `jiffies_to_timespec`, `jiffies_to_timeval`, ...

Time since Bootup (cont.)

- ★ Platform specific “Time Stamp Counter”
 - ▶ On x86
 - Header: `<asm/msr.h>`
 - API: `rdtsc(ul low_tsc_ticks, ul high_tsc_ticks);`
 - ▶ Getting it generically
 - Header: `<linux/timex.h>`
 - API: `read_current_timer(unsigned long *timer_val);`

Absolute Time

★ Header: `<linux/time.h>`

★ APIs

- ▶ `mktime(y, m, d, h, m, s)` – Seconds since Epoch
- ▶ `void do_gettimeofday(struct timeval *tv);`
- ▶ `struct timespec current_kernel_time(void);`

Delays

Long Delays

- ★ Busy wait: `cpu_relax`

```
while (time_before(jiffies, j1))  
    cpu_relax();
```

- ★ Yielding: `schedule/schedule_timeout`

```
while (time_before(jiffies, j1))  
    schedule();
```


Short Delays but Busy Waiting

- ★ Header: `<linux/delay.h>`
- ★ Arch. specific Header: `<asm/delay.h>`
- ★ APIs
 - `void ndelay(unsigned long ndelays);`
 - `void udelay(unsigned long udelays);`
 - `void mdelay(unsigned long mdelays);`

Long Delays: Back to Yielding

★ Header: `<linux/delay.h>`

★ APIs

- `void msleep(unsigned int millisecs);`
- `unsigned long msleep_interruptible(unsigned int millisecs);`
- `void ssleep(unsigned int secs);`

Timers

Kernel Timers

- ★ Back end of the various delays
- ★ Header: `<linux/timer.h>`
- ★ Type: `struct timer_list`
- ★ APIs
 - `void init_timer(struct timer_list *); /* Nullifies */`
 - `struct timer_list TIMER_INITIALIZER(f, t, p);`
 - `void add_timer(struct timer_list *);`
 - `void del_timer(struct timer_list *);`
 - `int mod_timer(struct timer_list *, unsigned long);`
 - `int del_timer_sync(struct timer_list *);`

Tasklets

- ★ Timers without specific Timing
- ★ Header: `<linux/interrupt.h>`
- ★ Type: `struct tasklet_struct`
- ★ APIs
 - `void tasklet_init(struct tasklet_struct *t, void (*func)(unsigned long), unsigned long data);`
 - `void tasklet_kill(struct tasklet_struct *t);`
 - `DECLARE_TASKLET(name, func, data);`
 - `tasklet_enable(t), tasklet_disable(t)`
 - `tasklet_[hi_]schedule(t);`

Work Queues

★ In context of “Special Kernel Thread”

★ Header: `<linux/workqueue.h>`

★ Types: `struct workqueue_struct`, `struct work_struct`

★ Work Queue APIs

- `q = create_workqueue(name);`
- `q = create_singlethread_workqueue(name);`
- `flush_workqueue(q);`
- `destroy_workqueue(q);`

★ Work APIs

- `DECLARE_WORK(w, void (*function)(void *), void *data);`
- `INIT_WORK(w, void (*function)(void *), void *data);`

★ Combined APIs

- `int queue_work(q, &w);`
- `int queue_delayed_work(q, &w, d);`
- `int cancel_delayed_work(&w);`

★ Global Shared Work Queue API

- `schedule_work(&w);`

Helper Interfaces

Other Helper Interfaces in Latest Kernels

- ★ User Mode Helper
- ★ Linked Lists
- ★ Hash Lists
- ★ Notifier Chains
- ★ Completion Interface
- ★ Kthread Helpers

What to Expect?

- ★ How to do programming in “Kernel C” for
 - Achieving Concurrency
 - With & without Locking
 - Wait Queues
 - Keeping Time
 - Relative & Absolute
 - Providing Delays
 - Long and Short
 - Busy Wait and Yielding
 - Timer Control
 - Kernel Timers
 - Tasklets
 - Work Queues

Any Queries?