

# Network Drivers

# What to Expect?

- ★ Understanding Network Subsystem
- ★ Network Drivers: A different category
- ★ Writing Network Drivers

# Network Subsystem

## ★ (Network) Protocol Stack

- Typically, the TCP/IP Stack
- Interfaces with the User Space through network interfaces, instead of device files
- Unlike other drivers, does not provide any /sys or /dev entries
- Interface Examples: eth0, wlan1, ...
- Resides in the <kernel\_source>/net folder

## ★ Network Interface Card (NIC) Drivers

- Driver for the Physical Network Cards
- Provides uniform hardware independent interface for the network protocol layer to access the card
- Typically, resides in the <kernel\_source>/drivers/net folder



# Related Data Structures

- ★ Writing a NIC or Network Driver involves
  - Interacting with the underlying Card over its I/O Bus
  - Providing the standard APIs to the Protocol Stack
- ★ This needs three kind of Data Structures
  - Core of Protocol Stack
    - struct sk\_buff
  - NIC & Protocol Stack Interface
    - struct net\_device
  - NIC I/O bus, e.g. PCI, USB, ...
    - I/O bus specific data structure

# struct sk\_buff

- ★ Network Packet Descriptor
- ★ Header: <linux/skbuff.h>
- ★ Driver relevant Fields
  - head – points to the start of the packet
  - tail – points to the end of the packet
  - data – points to the start of the pkt payload
  - end – points to the end of the packet payload
  - len – amount of data that the packet contains

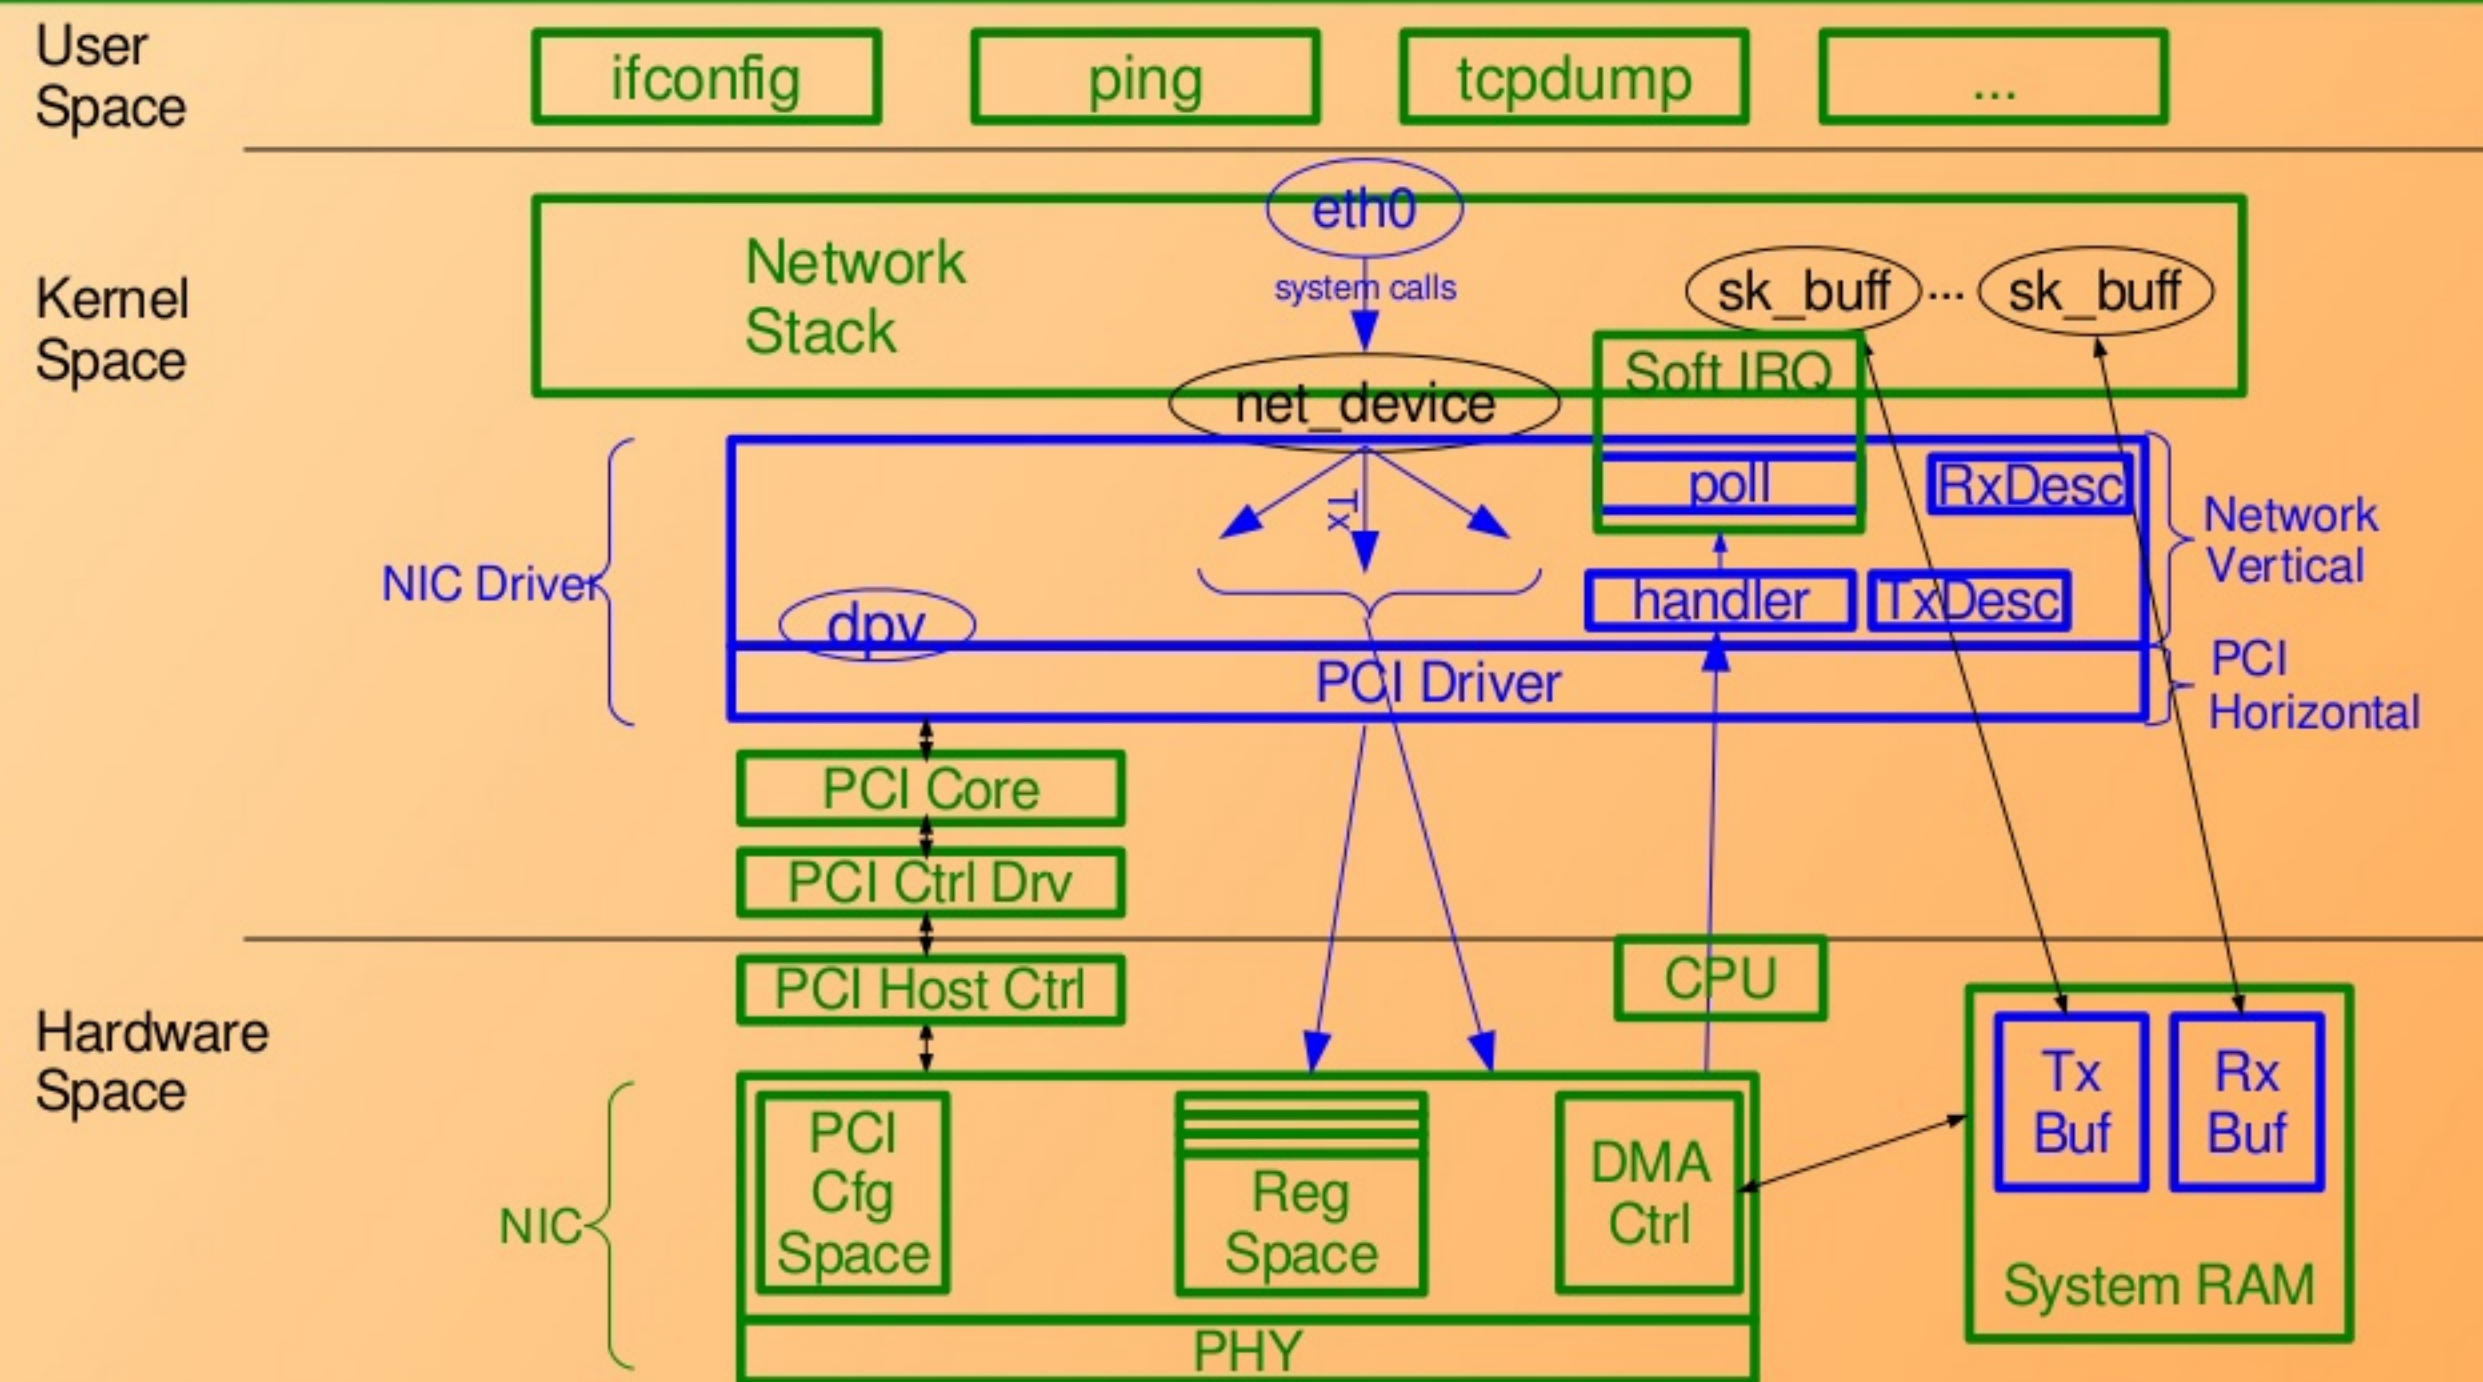


# Socket Buffer APIs

- ★ Header: <linux/skbuff.h>
- ★ SK Buffer Storage
  - `struct sk_buff *dev_alloc_skb(len);`
  - `dev_kfree_skb(skb);`
- ★ SK Buffer Operations
  - `void skb_reserve(struct sk_buff *, int len);`
  - `struct sk_buff *skb_clone(struct sk_buff *, gfp_t);`
  - `unsigned char *skb_put(struct sk_buff *, int len);`

# The Big Picture

(with a PCI NIC)





# struct net\_device

- ★ Header: <linux/netdevice.h>
- ★ Driver relevant Operation Fields
  - Activation: open, stop, ioctl
  - Data Transfer: start\_xmit, poll (NAPI: new API)
  - WatchDog: tx\_timeout, int watchdog\_timeo
  - Statistics: get\_stats, get\_wireless\_stats
    - Typically uses struct net\_device\_stats, populated earlier
  - Configuration: struct \*ethtool\_ops, change\_mtu
    - Structure defined in Header: <linux/ethtool.h>
  - Bus Specific: mem\_start, mem\_end
- ★ Latest kernels have all the operations moved under
  - struct net\_dev\_ops
  - And typically, a prefix ndo\_ added and some name changes



# Net Device Registration

★ Header: <linux/netdevice.h>

★ Net Device Storage

- `struct net_device *alloc_etherdev(sizeof_priv);`
- `struct net_device *alloc_ieee80211dev(sizeof_priv);`
- `struct net_device *alloc_irdadev(sizeof_priv);`
- `struct net_device *alloc_netdev(sizeof_priv, name, setup_fn);`
- `void free_netdev(struct net_device *);`

★ Registering the Net Device

- `int register_netdev(struct net_device *);`
- `void unregister_netdev(struct net_device *);`



# Network Device Open & Close

## ★ A typical Network Device Open

- Allocates ring buffers and associated sk\_buffs
- Initializes the Device
- Gets the MAC, ... from device EEPROM
- Requests firmware download, if needed
  - `int request_firmware(fw, name, device);`
- Register the interrupt handler(s)

## ★ A typical Network Device Close

- Would do the reverse in chronology reverse order



# Packet Receive Interrupt Handler

- ★ A typical receive interrupt handler
  - Minimally handles the packet received
    - Sanity checks
  - Puts back equal number of sk\_buffs for re-use
  - Passes the associated sk\_buffs (& ring buffers) to the protocol layer by the NET\_RX\_SOFTIRQ
  - On higher load, switch to poll mode, if supported, which then passes the associated sk\_buffs (& ring buffers) to the protocol layer by the NET\_RX\_SOFTIRQ

# Flow Control related APIs

- ★ For interaction with the protocol layer
- ★ Header: `<linux/netdevice.h>`
- ★ APIs
  - `void netif_start_queue(struct net_device *);`
  - `void netif_stop_queue(struct net_device *);`
  - `void netif_wake_queue(struct net_device *);`
  - `int netif_queue_stopped(struct net_device *);`



# Network Driver Examples

- ★ Driver: snull
- ★ Browse & Discuss
- ★ Driver for Realtek NIC 8136
- ★ Browse & Hack

# What all have we learnt?

- ☆ Linux Network Subsystem
- ☆ How are Network Drivers different?
- ☆ Writing Network Drivers
  - Key Data Structures & their APIs
    - sk\_buff & net\_device
  - Network Device Registration
  - Network Device Operations
  - Interrupt Handling for Packets received
  - Flow Control related APIs



Any Queries?