

System Calls

What to Expect?

- ★ W's of System Calls
- ★ System Call vs Library Function
- ★ System Call Tracing
- ★ Hands-On

W's of System Calls

- ★ User programs vs Kernel programs
 - Runs in different spaces
 - Runs with different privileges
 - User space not allowed access to Kernel space
 - But they need the Kernel services
- ★ OS provides service points
 - For User programs
 - To request services from the Kernel
- ★ In Linux, these are called System Calls

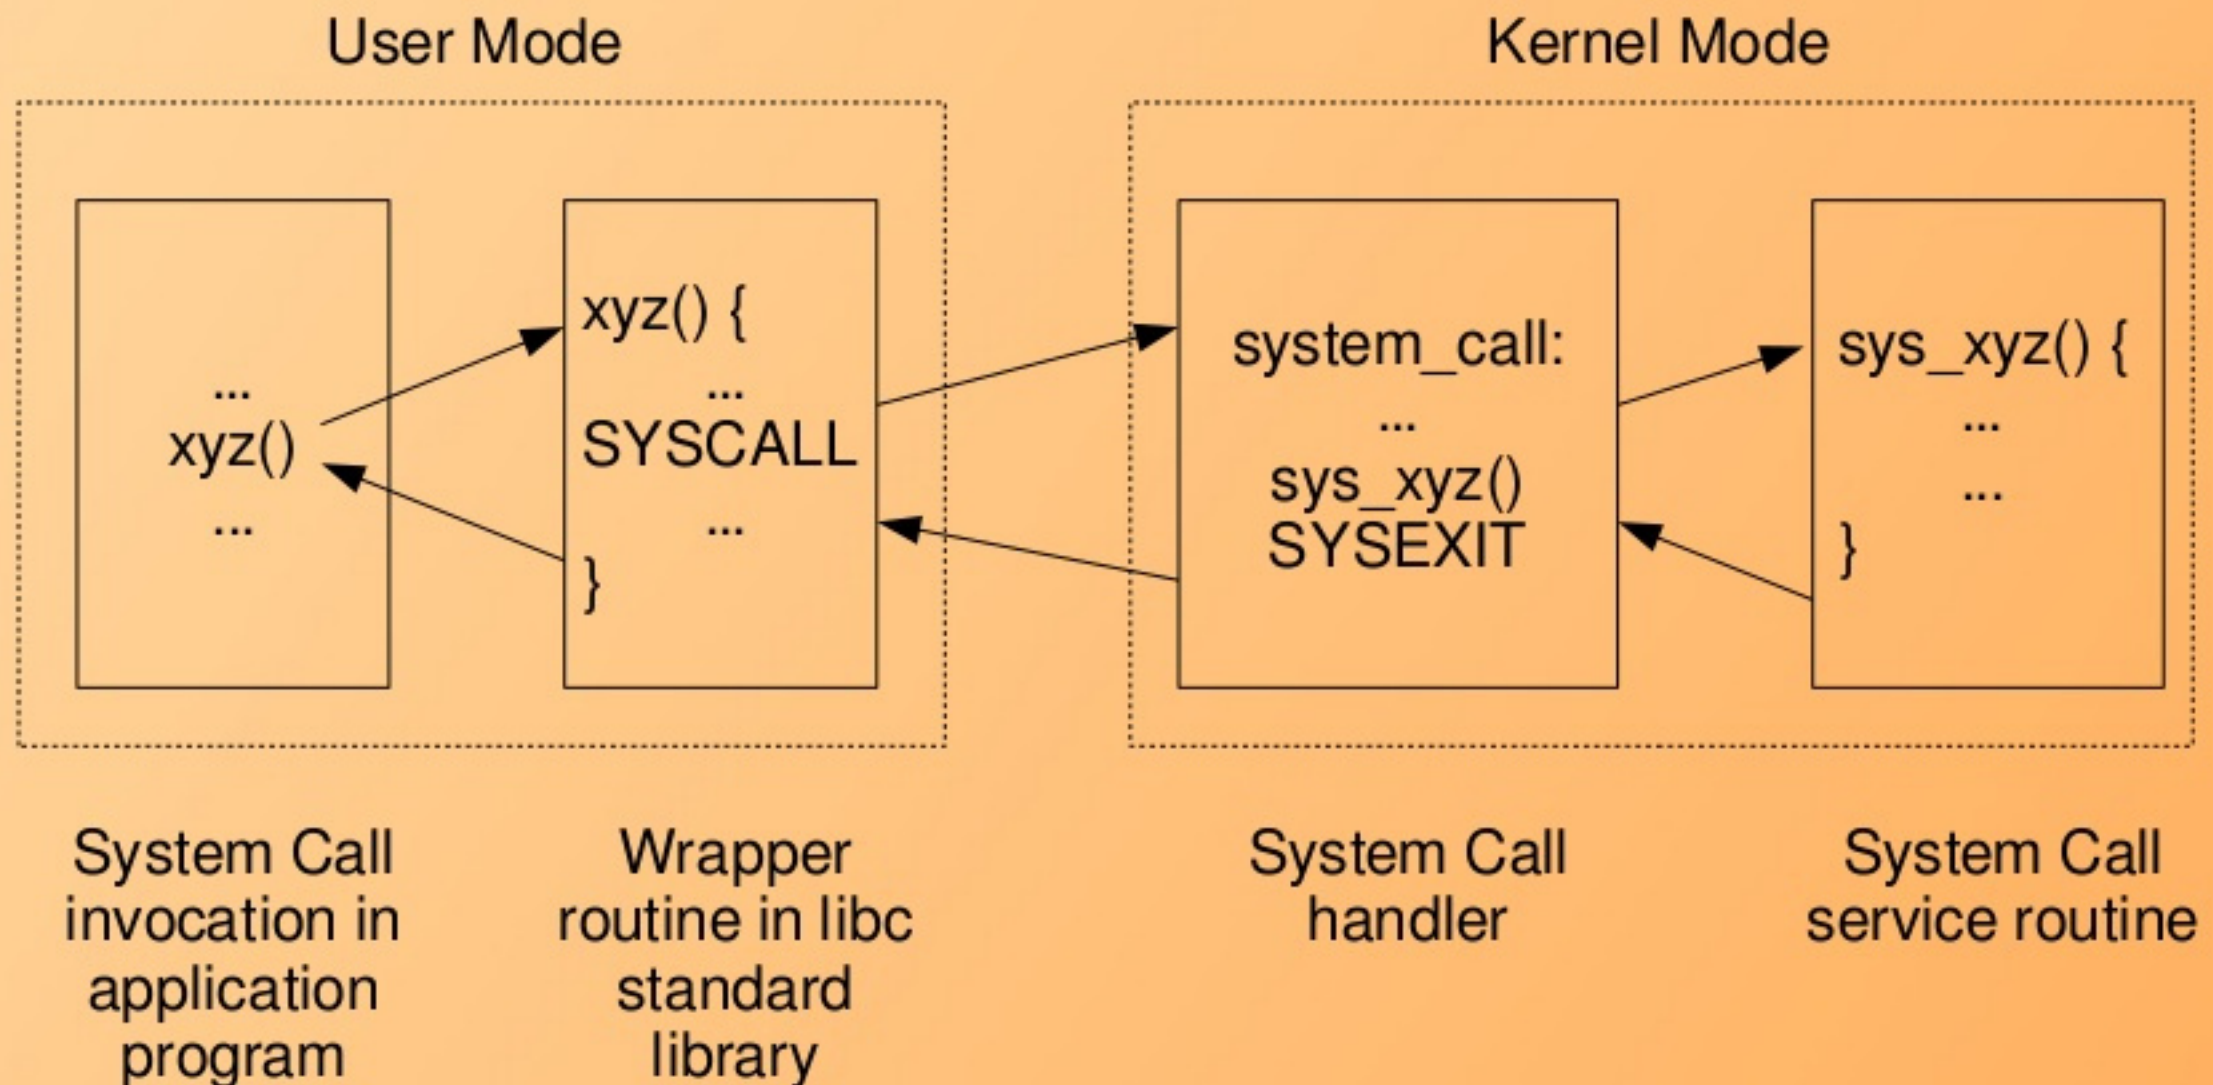
System Calls in Linux

- ★ About 300 in count
- ★ Listing: `/usr/include/asm/unistd.h`
- ★ Provide layer between
 - Kernel Space (typically hardware)
 - User Space (typically user process)
- ★ Serve three purposes
 - Provide an Abstracted h/w interface for user space
 - Ensures System security and stability
 - Makes Process Management easier

Working of a Linux System Call

- ★ Implemented as an ordinary function in the Linux Kernel
- ★ Executes like others in the Kernel Space
- ★ However, the call to that function isn't ordinary
- ★ When a user program makes a system call
 - Arguments are packaged up and handed to the kernel
 - A special procedure is required to transfer control to the kernel
 - Kernel takes over execution of the program until the call completes
 - Kernel transfers control back to the program with return value
- ★ Special procedure is typically achieved using “trap”

System Call Execution Flow



Linux System Call Wrappers

- ★ Every System Call has standard steps
- ★ GNU C library (glibc) abstracts them
 - By wrapping with functions of same name
 - For easy invocation
- ★ Examples
 - I/O functions: open, read, ...
- ★ We rarely invoke direct system calls
 - But rather these system call (wrapper) functions
- ★ Any Exception?
 - Custom defined system call – using `syscall(sno, ...)`

Contrast with a Library Function

- ★ A library function is an ordinary function
- ★ It resides in a library external to the program
 - But in the User Space only
- ★ Moreover, the call to it is also ordinary
 - Arguments placed in processor registers or the stack
 - Execution transferred to the start of the function
 - Typically resides in a loaded shared library
 - In the User Space only
- ★ Examples
 - fopen, printf, getopt, mkstemp (all from glibc)

Return Values

- ★ Library functions often return pointers
 - Example: `FILE * fp = fopen("harry","r");`
 - NULL indicates failure
- ★ System calls usually return an integer
 - Example: `int res = open("harry", O_RDONLY);`
 - Return value
 - `>= 0` indicates success
 - `< 0`, typically `-1` indicates failure, and error is set in `errno`
- ★ Note the counter intuitive return of System Calls
 - Opposite way round
 - Cannot use as Boolean

More Information

★ Manual Sections

- 2 System calls e.g. `_exit`, `read`, `write`
- 3 Library calls e.g. `exit`, `printf`
- 7 Miscellaneous e.g. `ascii`, `fifo`, `pthread`s
- 9 POSIX Programmer Manual

★ Info pages are also available

Tracing System Calls

- ★ Command: `strace <program> [args]`
- ★ Traces the execution of `<program>`
- ★ And Lists
 - ◆ System Calls made by `<program>`
 - ◆ Signals received by `<program>`
- ★ Controlled by various options
 - ◆ An interesting one is “-e”
- ★ Example
 - ◆ `strace cat /dev/null`

Pros & Cons

★ Pros

- ▶ System calls provide direct & hence more control over the kernel services
- ▶ Library functions abstract the nitty-gritty of architecture or OS specific details of the system calls
- ▶ Library functions can provide wrappers over repeated set of system calls

★ Cons

- ▶ Library functions may have overheads
- ▶ System calls at times may expose the underlying system dependency

Let's try some Examples

- ★ System Call Invocation
- ★ System calls vs Library functions
 - File Operations
- ★ Observe the various system calls invoked
 - Use strace

System Call: access

- ★ Helps to determine whether the calling process has access permission for a particular file
- ★ `int access(const char *pathname, int mode);`
 - ▶ `pathname` – Path to the file
 - ▶ `mode` – Accessibility checks. `F_OK` or a mask consisting of bitwise or of `R_OK`, `W_OK` and `X_OK`

System Call: fcntl

- ☆ Is used for performing the various advanced operations on an open file descriptor
- ☆ It allows to place a read or write lock on the file
- ☆ More than one process can hold the read lock, while only one process can hold the write lock
- ☆ `int fcntl(int fd, int cmd, ... /* arg */)`
 - fd – File descriptor
 - cmd – Operation to perform
 - ... – Arguments

What all have we learnt?

- ★ W's of System Calls
 - ▶ Working of a System Call & syscall()
 - ▶ System Call Wrapper Functions
- ★ System Call vs Library Function
 - ▶ Pros & Cons
- ★ System Call Tracing
- ★ Hands-On

Any Queries?